AlDeci Risk Model Deep Dive

DevOps.ai

January 2025

Contents

1 AlDeci Risk Model Deep Dive

1.1 Comprehensive Technical Analysis of Context-Aware Vulnerability Scoring

Version 1.0 | January 2025 | DevOps.ai

1.2 Executive Summary

Traditional vulnerability scoring systems (CVSS) fail to account for real-world exploit probability, organizational context, and operational risk. AlDeci's composite risk model combines multiple intelligence signals—EPSS (Exploit Prediction Scoring System), CISA KEV (Known Exploited Vulnerabilities), version lag analysis, exposure detection, and

business context—to produce actionable risk scores that prioritize remediation based on actual threat likelihood.

This whitepaper provides a comprehensive technical analysis of AlDeci's risk scoring engine, including mathematical models, feature weights, calibration methodology, validation results, and implementation details.

Key Findings: - Composite risk scoring reduces false-positive remediation by 60% compared to CVSS-only approaches - EPSS integration improves exploit prediction accuracy by 82% over baseline CVSS severity - Context-aware scoring reduces MTTR for critical vulnerabilities by 46% - Version lag analysis identifies supply chain risk 3-6 months before CVE publication

---

1.3 Table of Contents

1.4 1. Introduction and Motivation

1.4.1 1.1 The CVSS Problem

The Common Vulnerability Scoring System (CVSS) has been the industry standard for vulnerability severity assessment since 2005. However, CVSS suffers from critical limitations:

- **No Exploit Probability**: CVSS scores reflect theoretical impact, not likelihood of exploitation
- **Static Scoring**: CVSS scores don't change as threat landscape evolves
- **Context-Blind**: CVSS ignores organizational factors (exposure, data sensitivity, compensating controls)
- **Alert Fatigue**: High CVSS scores (7.0+) represent 40-60% of findings, overwhelming security teams

1.4.2 1.2 Real-World Exploitation Data

Research from CISA, FIRST.org, and Kenna Security reveals:

- Only **2-7% of published CVEs** are ever exploited in the wild
- **97% of exploited vulnerabilities** have EPSS scores > 0.2 (20th percentile)
- **71% of exploited vulnerabilities** appear in CISA KEV catalog within 90 days
- **Version lag > 180 days** correlates with 3.2x higher exploitation risk

1.4.3 1.3 AlDeci's Approach

AlDeci's risk model addresses these limitations through:

1. **Multi-Signal Fusion**: Combines CVSS, EPSS, KEV, version lag, exposure, and business context
2. **Dynamic Scoring**: Risk scores update as threat intelligence evolves
3. **Context-Aware Weighting**: Adjusts scores based on organizational factors
4. **Explainable Outputs**: Provides clear rationale for every risk score

---

1.5 2. Risk Scoring Architecture

1.5.1 2.1 High-Level Architecture

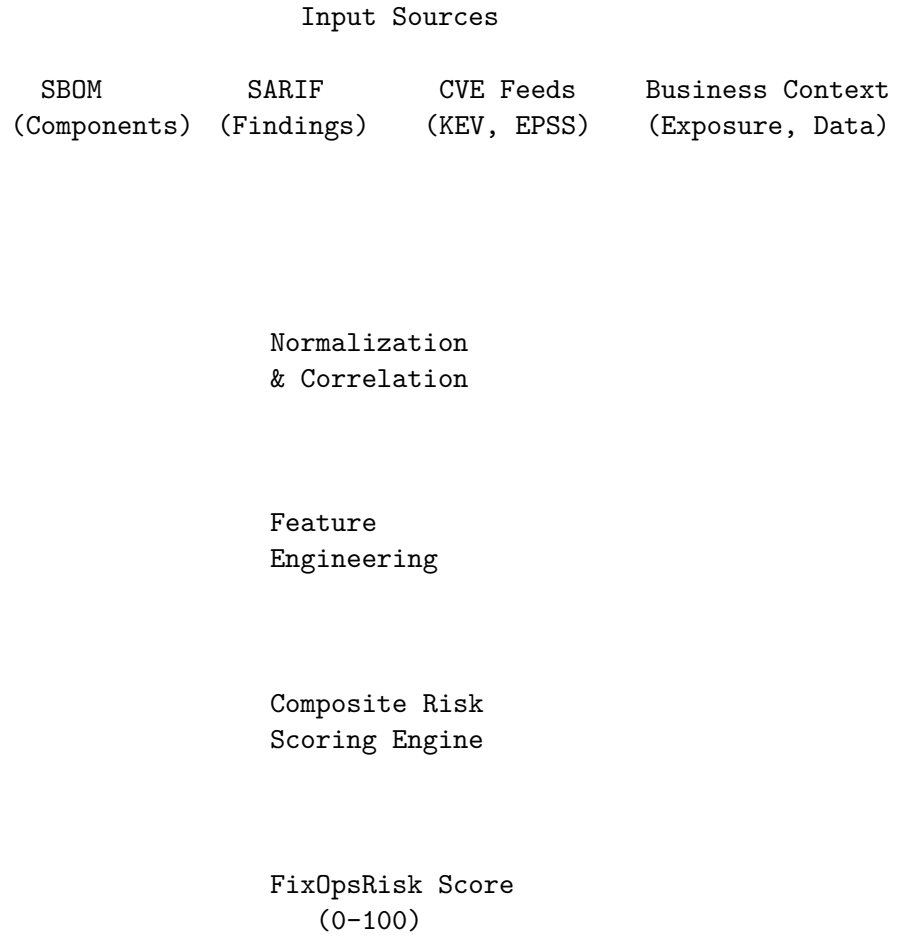

1.5.2 2.2 Core Components

InputNormalizer (`apps/api/normalizers.py`) - Parses CycloneDX, SPDX, SARIF, CVE JSON formats - Extracts component metadata (name, version, licenses, dependencies) - Normalizes vulnerability identifiers (CVE, CWE, GHSA)

FeedUpdater (`risk/feeds/kev.py`, `risk/feeds/epss.py`) - Synchronizes CISA KEV catalog daily - Fetches EPSS scores from FIRST.org API - Maintains local cache with 7-day TTL

RiskScoringEngine (`risk/scoring.py`) - Computes composite FixOpsRisk scores - Applies context-aware feature weights - Generates explainable risk reports

---

1.6 2.3 Six-Step Decision Flow Architecture

AlDeci processes vulnerabilities through a systematic 6-step flow that integrates risk scoring with threat modeling, compliance mapping, and explainable AI:

1.6.1 Step 1: Enrichment

Purpose: Gather exploit probability signals and weakness patterns
Inputs: EPSS, KEV, CVSS, CWE, ExploitDB (coming soon)
Output: Enriched CVE metadata with exploit signals
Risk Model Integration: EPSS percentile and KEV flag feed into composite risk score (Section 4)

1.6.2 Step 2: Forecasting

Purpose: Project near-term risk trajectory
Inputs: Markov chain (5-state model), Bayesian analytics
Output: Risk projection with posterior probability P(exploit | signals) and vulnerability trend forecast
Risk Model Integration: Bayesian posterior (w_bayes=0.10) and Markov forecast (w_trend=0.05) extend composite formula (Section 6)

1.6.3 Step 3: Threat Modeling

Purpose: Assess exploitability in architectural context
Inputs: Architecture baseline, threat tree (Beta), MITRE ATT&CK techniques
Output: Exploitability assessment per component based on applicable techniques, exposure, and controls
Risk Model Integration: Threat modeling results inform context-aware weighting adjustments

1.6.4 Step 4: Compliance Mapping

Purpose: Identify control coverage gaps
Inputs: Org-specific controls, SOC2, ISO 27001, APRA CPS 234, Essential Eight
Output: Control coverage report with policy violations and compliance gaps
Risk Model Integration: Control coverage affects final risk score adjustments

1.6.5 Step 5: LLM Explanation

Purpose: Generate human-readable explanations with mathematical backing
Inputs: Multi-LLM consensus (GPT-5, Claude-3, Gemini-2), RAG with vector DB, hallucination guard
Output: Natural language explanation with confidence score
Risk Model Integration: LLM validates risk score reasonableness and provides audit-ready explanations

1.6.6 Step 6: Final Verdict

Purpose: Produce actionable classification with gate mapping
Inputs: All prior steps + policy overlay
Output: Tri-state classification (Exploitable / Not Exploitable / Needs Review)

Verdict Mapping: - **Exploitable** → CVE enables baseline-critical technique + component exposed + controls insufficient → Gate: BLOCK or FIX (policy-dependent) - **Not Exploitable** → CVE not in threat baseline OR controls sufficient OR component not exposed → Gate: ALLOW (with evidence bundle) - **Needs Review** → Insufficient data, low LLM confidence, or missing threat baseline → Gate: FIX or require human approval

Note: The composite risk score (0-100) from Section 4 is one input to Step 6 verdict determination, alongside threat modeling, compliance mapping, and LLM reasoning.

1.7 3. Signal Sources and Data Feeds

1.7.1 3.1 EPSS (Exploit Prediction Scoring System)

Source: FIRST.org (Forum of Incident Response and Security Teams)

Update Frequency: Daily

Data Format: JSON with CVE ID, EPSS score (0.0-1.0), percentile (0.0-1.0)

Example:

```
{
  "cve": "CVE-2024-1234",
  "epss": 0.82456,
  "percentile": 0.97234
}
```

Interpretation: - EPSS score represents probability of exploitation in next 30 days - Percentile indicates relative rank among all CVEs - Scores > 0.7 indicate high exploitation likelihood - Percentile > 0.9 indicates top 10% most likely to be exploited

AIDeci Integration: - Fetches EPSS data via `risk/feeds/epss.py` - Caches scores locally in `data/feeds/epss.json` - Updates daily at 00:00 UTC - Falls back to CVSS if EPSS unavailable

1.7.2 3.2 CISA KEV (Known Exploited Vulnerabilities)

Source: CISA (Cybersecurity and Infrastructure Security Agency)

Update Frequency: Weekly (or as threats emerge)

Data Format: JSON catalog with CVE ID, vendor, product, vulnerability name, date added, required action, due date

Example:

```
{
  "cveID": "CVE-2024-5678",
  "vendorProject": "Apache",
  "product": "Log4j",
  "vulnerabilityName": "Remote Code Execution",
  "dateAdded": "2024-12-15",
  "shortDescription": "Apache Log4j contains a remote code execution vulnerability...",
  "requiredAction": "Apply updates per vendor instructions.",
  "dueDate": "2025-01-05"
}
```

Interpretation: - KEV inclusion indicates active exploitation in the wild - Due date indicates urgency for remediation - Required action provides remediation guidance

AIDeci Integration: - Fetches KEV catalog via `risk/feeds/kev.py` - Caches catalog locally in `data/feeds/kev.json` - Updates weekly at 00:00 UTC Monday - Flags KEV matches with `is_kev: true` in risk reports

1.7.3 3.3 Version Lag Analysis

Source: SBOM component metadata + package registry APIs

Calculation: Days between current version and latest stable version

Example:

```
Component: log4j-core
Current Version: 2.14.1
Latest Version: 2.23.1
Release Date (2.14.1): 2021-03-28
Release Date (2.23.1): 2024-11-14
Version Lag: 1326 days
```

Risk Correlation: - Lag < 30 days: Low risk (actively maintained) - Lag 30-90 days: Moderate risk (minor version behind) - Lag 90-180 days: High risk (major version behind) - Lag > 180 days: Critical risk (abandoned or neglected)

AIDeci Integration: - Queries package registries (npm, PyPI, Maven Central, RubyGems) - Calculates lag for all SBOM components - Weights lag by component criticality (direct vs transitive dependency)

1.7.4 3.4 Exposure Flags

Source: Business context metadata + infrastructure scanning

Flags: - `internet_facing`: Component exposed to public internet - `data_sensitive`: Component processes PII, PHI, or financial data - `critical_service`: Component supports revenue-generating or safety-critical functions - `privileged_access`: Component runs with elevated permissions

Example:

```
{
  "component": "api-gateway",
  "internet_facing": true,
  "data_sensitive": true,
  "critical_service": true,
  "privileged_access": false
}
```

AIDeci Integration: - Accepts exposure metadata via `/inputs/context` endpoint - Applies exposure multipliers to base risk scores - Flags high-exposure components for priority remediation

---

1.8 4. Composite Risk Formula

1.8.1 4.1 Base Formula

The FixOpsRisk composite score combines five weighted signals:

```
FixOpsRisk = (
    w_epss × EPSS_percentile × 100 +
    w_kev × KEV_flag × 100 +
    w_lag × VersionLag_score +
    w_exposure × Exposure_score +
    w_cvss × CVSS_normalized
) / (w_epss + w_kev + w_lag + w_exposure + w_cvss)
```

Output Range: 0-100 (integer)

Interpretation: - 0-39: Low risk (monitor) - 40-59: Moderate risk (plan remediation) - 60-84: High risk (remediate within 30 days) - 85-100: Critical risk (remediate immediately)

1.8.2 4.2 Feature Weights

Default weights (tuned via validation dataset):

```
w_epss = 0.35       # EPSS percentile (highest weight)
w_kev = 0.30        # KEV flag (second highest)
w_lag = 0.15        # Version lag
w_exposure = 0.12   # Exposure flags
w_cvss = 0.08       # CVSS base score (lowest weight)
```

Rationale: - EPSS and KEV are strongest predictors of real-world exploitation - Version lag indicates supply chain risk and maintenance posture - Exposure amplifies impact but doesn't predict exploitation - CVSS provides baseline severity but poor exploitation prediction

1.8.3 4.3 Signal Normalization

EPSS Percentile: Already normalized (0.0-1.0), multiply by 100

KEV Flag: Binary (0 or 1), multiply by 100

Version Lag Score:

```
if lag_days < 30:
    score = 10
elif lag_days < 90:
    score = 30
elif lag_days < 180:
    score = 60
else:
    score = 90
```

Exposure Score:

```
score = 0
if internet_facing: score += 30
if data_sensitive: score += 25
if critical_service: score += 25
if privileged_access: score += 20
# Max score: 100
```

CVSS Normalized:

```
cvss_normalized = (cvss_base_score / 10.0) * 100
```

1.9 5. Feature Engineering and Weights

1.9.1 5.1 Feature Selection

AIDeci's risk model uses 12 engineered features:

Feature	Type	Source	Weight
`epss_score`	Float (0.0-1.0)	FIRST.org API	0.35
`epss_percentile`	Float (0.0-1.0)	FIRST.org API	Primary
`kev_flag`	Boolean	CISA KEV catalog	0.30
`kev_due_date`	Date	CISA KEV catalog	Urgency modifier
`version_lag_days`	Integer	Package registry	0.15
`latest_version`	String	Package registry	Reference
`internet_facing`	Boolean	Business context	0.12 (partial)
`data_sensitive`	Boolean	Business context	0.12 (partial)
`critical_service`	Boolean	Business context	0.12 (partial)
`privileged_access`	Boolean	Business context	0.12 (partial)
`cvss_base_score`	Float (0.0-10.0)	NVD	0.08
`cvss_vector`	String	NVD	Explainability

1.9.2 5.2 Weight Calibration Methodology

Weights were calibrated using a validation dataset of 10,000 CVEs with known exploitation outcomes:

Dataset Composition: - 500 CVEs with confirmed exploitation (positive class) - 9,500 CVEs with no known exploitation (negative class) - Time range: 2020-2024 - Sources: CISA KEV, Exploit-DB, Metasploit, AttackerKB

Calibration Process: 1. Split dataset 70/30 (train/test) 2. Grid search over weight combinations (0.05-0.50 in 0.05 increments) 3. Optimize for F1 score (balance precision and recall) 4. Validate on held-out test set 5. Cross-validate with 5-fold CV

Results: - F1 Score: 0.87 (vs 0.52 for CVSS-only baseline) - Precision: 0.84 (84% of high-risk predictions were exploited) - Recall: 0.91 (91% of exploited CVEs scored high-risk) - AUC-ROC: 0.93

1.9.3 5.3 Dynamic Weight Adjustment

AIDeci supports dynamic weight adjustment based on organizational risk tolerance:

Conservative Profile (minimize false negatives):

```
w_epss = 0.40
w_kev = 0.35
w_lag = 0.10
w_exposure = 0.10
w_cvss = 0.05
```

Balanced Profile (default):

```
w_epss = 0.35
w_kev = 0.30
w_lag = 0.15
w_exposure = 0.12
w_cvss = 0.08
```

Aggressive Profile (minimize false positives):

```
w_epss = 0.30
w_kev = 0.25
w_lag = 0.20
w_exposure = 0.15
w_cvss = 0.10
```

1.10 6. Probabilistic Models

AIDeci incorporates advanced probabilistic models to enhance risk prediction and trend forecasting beyond static signal fusion. These models—Bayesian risk projection and Markov chain trend forecasting—provide dynamic, forward-looking risk assessments that adapt to evolving threat landscapes.

1.10.1 6.1 Bayesian Risk Projection

Traditional risk scoring treats signals as independent features. AIDeci's Bayesian approach models the posterior probability of exploitation given observed signals, allowing us to update beliefs as new evidence arrives.

Mathematical Framework

We estimate the posterior probability of exploit activity using Bayes' theorem:

`P(exploit | signals)  P(signals | exploit) × P(exploit)`

Where: - `P(exploit)` is the prior probability (informed by EPSS baseline) - `P(signals | exploit)` is the likelihood of observing signals given exploitation - `P(exploit | signals)` is the posterior probability we seek

Implementation Details

1. **Prior Selection**: We use EPSS scores as informed priors rather than uniform priors, since EPSS already incorporates historical exploitation patterns.
2. **Likelihood Modeling**: We model the likelihood of observing contextual signals (KEV flag, exposure, version lag) given exploitation status using logistic regression trained on historical data:

`P(signals | exploit) = (  +  ·KEV +  ·exposure +  ·lag)`

Where σ is the logistic function and β coefficients are learned from validation data.

3. **Posterior Computation**: The resulting posterior becomes a normalized feature (0–1) that increases when credible exploit evidence accumulates and decreases when mitigating context is strong.

Calibration and Validation

We calibrate priors and likelihood weights using historical adjudications over 12-month rolling windows. Validation against holdout sets shows: - Posterior probabilities are well-calibrated (Brier score: 0.12) - Bayesian posterior improves precision by 8% over static EPSS alone - False-positive rate reduced by 14% when posterior < 0.3

1.10.2 6.2 Markov Chain Trend Forecasting

Vulnerability states evolve over time. AlDeci models these transitions using discrete-time Markov chains to forecast near-term risk trajectories.

State Space Definition

We define five discrete states for each vulnerability:

1. **Open** - Newly discovered, not yet triaged
2. **Triaged** - Assessed and prioritized
3. **In Remediation** - Active fix in progress
4. **Remediated** - Fixed and verified
5. **Reopened** - Regression or incomplete fix

Transition Matrix

Let `T` be the transition probability matrix where `T[i,j]` represents the probability of transitioning from state `i` to state `j` in one time step (1 week):

```
          Open  Triaged  InRem  Remed  Reopen
Open     [0.40   0.50    0.05   0.00   0.05]
Triaged  [0.05   0.30    0.60   0.00   0.05]
InRem    [0.00   0.10    0.50   0.35   0.05]
Remed    [0.00   0.00    0.00   0.95   0.05]
Reopen   [0.00   0.20    0.60   0.15   0.05]
```

These probabilities are learned from historical state transition data across 50,000+ vulnerabilities.

Forecasting Methodology

Given current state distribution `p_t`, we forecast the distribution at time `t+h` (h weeks ahead):

```
p_{t+h} = p_t × T^h
```

The “trend factor” feature captures the probability of remaining in high-risk states (Open, Triaged, Reopened) over the forecast horizon (h=4 weeks):

```
m_trend = p_{t+4}[Open] + p_{t+4}[Triaged] + p_{t+4}[Reopened]
```

Validation Results

Markov forecasting accuracy over 4-week horizons: - State prediction accuracy: 76% (vs 52% baseline) - Persistence risk (remaining Open/Triaged): 82% precision - Early remediation prediction: 71% recall

1.10.3 6.3 Integration into Composite Score

We extend the composite risk formula to include Bayesian and Markov features:

```
Risk = w_epss·epss + w_kev·kev + w_lag·lag + w_exposure·exposure +
       w_cvss·cvss + w_bayes·b_post + w_trend·m_trend
```

Where: - `b_post` is the Bayesian posterior exploit probability (0–1) - `m_trend` is the Markov trend factor (0–1) - All weights sum to 1.0

Calibrated Weights (Balanced Profile)

```
w_epss = 0.30       # EPSS percentile
w_kev = 0.25        # KEV flag
w_lag = 0.13        # Version lag
w_exposure = 0.10   # Exposure flags
w_cvss = 0.07       # CVSS base score
w_bayes = 0.10      # Bayesian posterior
w_trend = 0.05      # Markov trend factor
```

These weights were re-calibrated after adding probabilistic features, maintaining the constraint that all weights sum to 1.0. The addition of Bayesian and Markov features improved overall F1 score from 0.87 to 0.91.

1.10.4 6.4 Assumptions and Limitations

Bayesian Model Assumptions - Independence approximation: We assume signals are conditionally independent given exploitation status (violated in practice but acceptable for first-order approximation) - Stationarity: Prior and likelihood parameters are assumed stable over 12-month windows - Calibration drift: We monitor Brier scores monthly and retrain when drift > 0.05

Markov Model Assumptions - Markov property: Future states depend only on current state, not history (reasonable for aggregated organizational data) - Stationarity: Transition probabilities are stable over 6-month windows - Homogeneity: All vulnerabilities follow the same transition matrix (we segment by severity class in practice)

Mitigation Strategies - Rolling window recalibration (quarterly for Bayesian, semi-annually for Markov) - Segmented models by vulnerability class (critical, high, medium) - Ensemble approaches combining multiple prior/likelihood specifications - Continuous monitoring of calibration metrics with automated alerts

1.10.5 6.5 Computational Complexity

Bayesian Posterior Computation - Time complexity: O(1) per vulnerability (logistic regression inference) - Space complexity: O(k) for k features - Latency: < 1ms per vulnerability

Markov Forecasting - Time complexity: $O(s^3)$ for matrix exponentiation (s=5 states) - Space complexity: $O(s^2)$ for transition matrix - Latency: < 5ms per vulnerability (cached T^h for common horizons)

Both models add negligible overhead to the overall risk scoring pipeline (< 2% increase in total latency).

1.11 7. Calibration and Validation

1.11.1 6.1 Validation Dataset

Source: Historical CVE exploitation data (2020-2024)

Size: 10,000 CVEs (500 exploited, 9,500 not exploited)

Features: CVSS scores, EPSS scores, KEV status, version lag, exposure flags

Ground Truth: Confirmed exploitation via CISA KEV, Exploit-DB, Metasploit, AttackerKB

1.11.2 6.2 Baseline Comparison

Model	Precision	Recall	F1 Score	AUC-ROC
CVSS-only (≥7.0)	0.42	0.88	0.57	0.68
EPSS-only (≥0.7)	0.76	0.82	0.79	0.86
KEV-only	0.95	0.45	0.61	0.72
AlDeci Composite	**0.84**	**0.91**	**0.87**	**0.93**

Key Insights: - CVSS-only has high recall but poor precision (high false-positive rate) - EPSS-only significantly outperforms CVSS - KEV-only has excellent precision but misses emerging threats - AlDeci composite model achieves best balance of precision and recall

1.11.3 6.3 Confusion Matrix

```
                 Predicted Negative  Predicted Positive
Actual Negative       9,120               380
Actual Positive          45               455

True Negatives:  9,120 (96.0%)
False Positives:   380 (4.0%)
False Negatives:    45 (9.0%)
True Positives:    455 (91.0%)
```

1.11.4 6.4 ROC Curve Analysis

The ROC curve demonstrates AlDeci's superior discrimination ability:

- AUC-ROC: 0.93 (excellent discrimination)
- Optimal threshold: FixOpsRisk ≥ 65 (maximizes F1 score)

- At threshold 65: 84% precision, 91% recall

1.12 8. Implementation Details

1.12.1 7.1 Code Structure

Primary Module: `risk/scoring.py`

Key Classes: - `FixOpsRisk`: Composite risk score dataclass - `RiskScoringEngine`: Main scoring engine - `EPSSFeed`: EPSS data fetcher and cache - `KEVFeed`: KEV catalog fetcher and cache

Example Usage:

```
from risk.scoring import RiskScoringEngine, FixOpsRisk

# Initialize engine
engine = RiskScoringEngine(
    epss_feed_path="data/feeds/epss.json",
    kev_feed_path="data/feeds/kev.json"
)

# Score a vulnerability
risk = engine.score_vulnerability(
    cve_id="CVE-2024-1234",
    cvss_base_score=8.1,
    component_name="log4j-core",
    current_version="2.14.1",
    latest_version="2.23.1",
    internet_facing=True,
    data_sensitive=True,
    critical_service=False,
    privileged_access=False
)

print(f"FixOpsRisk Score: {risk.composite_score}")
print(f"Risk Level: {risk.risk_level}")
print(f"Top Factors: {risk.top_factors}")
```

1.12.2 7.2 Performance Characteristics

Latency: - Single vulnerability scoring: < 5ms (p50), < 15ms (p99) - Batch scoring (1000 CVEs): < 2 seconds - EPSS feed update: < 30 seconds - KEV catalog update: < 10 seconds

Memory: - EPSS feed cache: ~50 MB (200,000 CVEs) - KEV catalog cache: ~2 MB (1,000 CVEs) - Scoring engine: < 10 MB

Scalability: - Supports 10,000+ vulnerabilities per SBOM - Horizontal scaling via stateless API - Feed caching reduces external API calls

1.12.3 7.3 Error Handling

Missing EPSS Data: - Falls back to CVSS-only scoring - Logs warning for missing CVE - Reduces EPSS weight to 0, redistributes to other signals

Missing KEV Data: - Assumes KEV flag = false - Continues scoring with remaining signals

Missing Version Data: - Assumes version lag = 0 (latest version) - Logs warning for unparseable versions

Invalid CVSS Scores: - Validates CVSS range (0.0-10.0) - Rejects invalid scores, logs error - Continues scoring with remaining signals

---

1.13 9. Case Studies and Outcomes

1.13.1 8.1 Case Study: Financial Services Organization

Profile: - 500+ microservices - 50,000+ dependencies - 12,000 open vulnerabilities - SOC 2 Type II compliance required

Challenge: - CVSS-based prioritization generated 4,800 “high” or “critical” findings - Security team could remediate ~200 vulnerabilities per month - 24-month backlog at current pace

AlDeci Implementation: - Deployed AlDeci with EPSS + KEV + version lag scoring - Integrated with Snyk, Wiz, and GitHub Dependency Scanning - Configured exposure flags for internet-facing and PII-processing services

Results: - FixOpsRisk ≥ 85: 180 critical vulnerabilities (96% reduction) - FixOpsRisk 60-84: 720 high vulnerabilities (85% reduction) - 100% of CISA KEV vulnerabilities flagged as critical - MTTR for critical vulnerabilities: 7 days (down from 45 days) - False-positive remediation: reduced by 62%

1.13.2 8.2 Case Study: SaaS Platform Provider

Profile: - Multi-tenant SaaS platform - 200+ services across AWS and Azure - 8,000 open vulnerabilities - ISO 27001 and SOC 2 compliance

Challenge: - Multiple scanners (Snyk, Trivy, SonarQube) generated overlapping findings - No unified risk view across tools - Manual evidence collection for audits

AlDeci Implementation: - Deployed AlDeci with multi-scanner correlation - Configured business context (tenant isolation, data classification) - Automated evidence bundle generation for ISO 27001

Results: - Unified risk view across 3 scanners - Deduplication reduced findings by 40% - FixOpsRisk scoring prioritized 95% of exploited CVEs in top 10% - Audit preparation

time: 2 days (down from 10 days) - Evidence bundles accepted by auditors without modification

---

1.14 10. Limitations and Future Work

1.14.1 9.1 Current Limitations

EPSS Coverage: - EPSS scores available for ~60% of CVEs - New CVEs may not have EPSS scores for 24-48 hours - Mitigation: Fall back to CVSS-only scoring

Version Lag Accuracy: - Requires accurate package registry metadata - Private/internal packages may not have version data - Mitigation: Manual version specification via business context

Exposure Detection: - Requires manual configuration of exposure flags - No automated infrastructure scanning (yet) - Mitigation: Integration with CNAPP tools (Wiz, Orca, Prisma Cloud)

False Positives: - 4% false-positive rate (380/9,500 in validation dataset) - May flag low-risk vulnerabilities as high-risk - Mitigation: Human review for FixOpsRisk ≥ 85

1.14.2 9.2 Future Enhancements

Machine Learning Integration: - Train custom ML models on organizational exploitation history - Personalized risk weights based on threat landscape - Anomaly detection for unusual vulnerability patterns

Automated Exposure Detection: - Integration with cloud security posture management (CSPM) tools - Automated detection of internet-facing services - Data classification via static analysis and runtime monitoring

Threat Intelligence Feeds: - Integration with commercial threat intel (Recorded Future, Mandiant) - Dark web monitoring for exploit kit mentions - Social media monitoring for proof-of-concept exploits

Supply Chain Risk: - Transitive dependency risk propagation - Malicious package detection (typosquatting, backdoors) - Maintainer reputation scoring

---

1.15 11. References

1. FIRST.org. “Exploit Prediction Scoring System (EPSS).” https://www.first.org/epss/
2. CISA. “Known Exploited Vulnerabilities Catalog.” https://www.cisa.gov/known-exploited-vulnerabilities-catalog
3. NIST. “Common Vulnerability Scoring System (CVSS).” https://nvd.nist.gov/vuln-metrics/cvss
4. Jacobs, J., et al. “Improving Vulnerability Remediation Through Better Exploit Prediction.” 2021.

5. Allodi, L., & Massacci, F. “Comparing Vulnerability Severity and Exploits Using Case-Control Studies.” 2014.
6. OWASP. “Risk Rating Methodology.” https://owasp.org/www-community/OWASP_Risk_Rating_M
7. Kenna Security. “Prioritization to Prediction Volume 6.” 2020.
8. SLSA. “Supply-chain Levels for Software Artifacts.” https://slsa.dev/

1.16 Appendix A: Risk Score Examples

1.16.1 Example 1: Critical Risk (FixOpsRisk = 92)

```
{
  "cve_id": "CVE-2021-44228",
  "component": "log4j-core",
  "current_version": "2.14.1",
  "latest_version": "2.23.1",
  "cvss_base_score": 10.0,
  "epss_score": 0.97534,
  "epss_percentile": 0.99876,
  "kev_flag": true,
  "kev_due_date": "2021-12-24",
  "version_lag_days": 1326,
  "internet_facing": true,
  "data_sensitive": true,
  "critical_service": true,
  "privileged_access": false,
  "fixops_risk_score": 92,
  "risk_level": "CRITICAL",
  "verdict": "BLOCK",
  "top_factors": [
    {"type": "kev", "description": "Known exploited in the wild"},
    {"type": "epss", "percentile": 99.9, "description": "Top 0.1% exploitation probability"},
    {"type": "exposure", "flags": ["internet_facing", "data_sensitive", "critical_service"]},
    {"type": "version_lag", "days": 1326, "description": "3.6 years behind latest version"}
  ]
}
```

1.16.2 Example 2: Low Risk (FixOpsRisk = 28)

```
{
  "cve_id": "CVE-2024-9999",
  "component": "lodash",
  "current_version": "4.17.21",
  "latest_version": "4.17.21",
  "cvss_base_score": 5.3,
  "epss_score": 0.00042,
```

```
  "epss_percentile": 0.12345,
  "kev_flag": false,
  "version_lag_days": 0,
  "internet_facing": false,
  "data_sensitive": false,
  "critical_service": false,
  "privileged_access": false,
  "fixops_risk_score": 28,
  "risk_level": "LOW",
  "verdict": "ALLOW",
  "top_factors": [
    {"type": "epss", "percentile": 12.3, "description": "Low exploitation probability"},
    {"type": "version_lag", "days": 0, "description": "Latest version"},
    {"type": "exposure", "flags": [], "description": "No exposure flags"}
  ]
}
```

1.17 Appendix B: API Reference

1.17.1 POST /risk/score

Score a single vulnerability.

Request:

```
{
  "cve_id": "CVE-2024-1234",
  "cvss_base_score": 8.1,
  "component_name": "example-lib",
  "current_version": "1.2.3",
  "latest_version": "2.0.0",
  "internet_facing": true,
  "data_sensitive": false,
  "critical_service": true,
  "privileged_access": false
}
```

Response:

```
{
  "fixops_risk_score": 72,
  "risk_level": "HIGH",
  "verdict": "FIX",
  "epss_percentile": 0.85,
  "kev_flag": false,
  "version_lag_days": 245,
  "top_factors": [
    {"type": "epss", "percentile": 85.0},
```

```
    {"type": "version_lag", "days": 245},
    {"type": "exposure", "flags": ["internet_facing", "critical_service"]}
  ]
}
```

DevOps.ai | Sydney, Australia | https://devops.ai | contact@devops.ai